

Import Library

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

Load and Preprocess the Dataset

```
In [2]: # Step 1: Load and Preprocess the Dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Normalize pixel values

```
In [5]: # Normalize pixel values to range [0, 1]
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
```

Reshape data to fit the input requirements of a CNN

```
In [9]: # Reshape data to fit the input requirements of a CNN
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)
```

Convert labels to one-hot encoding

```
In [11]: # Convert labels to one-hot encoding
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)
```

Compile the Model

```
In [15]: # Step 3: Compile the Model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

Build the CNN Model

```
In [ ]: # Step 2: Build the CNN Model
model = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax')
])
```

Train the Model

```
In [17]: # Step 4: Train the Model
history = model.fit(x_train, y_train, validation_split=0.2, epochs=10, batch_size=128, verbose=1)
```

Epoch 1/10
 375/375 ————— 10s 20ms/step - accuracy: 0.7736 - loss: 0.7148 - val_accuracy: 0.9789 - val_loss: 0.0725
 Epoch 2/10
 375/375 ————— 7s 19ms/step - accuracy: 0.9684 - loss: 0.1154 - val_accuracy: 0.9772 - val_loss: 0.0683
 Epoch 3/10
 375/375 ————— 7s 18ms/step - accuracy: 0.9751 - loss: 0.0827 - val_accuracy: 0.9872 - val_loss: 0.0463
 Epoch 4/10
 375/375 ————— 7s 19ms/step - accuracy: 0.9810 - loss: 0.0595 - val_accuracy: 0.9868 - val_loss: 0.0485
 Epoch 5/10
 375/375 ————— 7s 18ms/step - accuracy: 0.9832 - loss: 0.0566 - val_accuracy: 0.9890 - val_loss: 0.0394
 Epoch 6/10
 375/375 ————— 7s 18ms/step - accuracy: 0.9865 - loss: 0.0444 - val_accuracy: 0.9901 - val_loss: 0.0366
 Epoch 7/10
 375/375 ————— 7s 18ms/step - accuracy: 0.9881 - loss: 0.0386 - val_accuracy: 0.9889 - val_loss: 0.0392
 Epoch 8/10
 375/375 ————— 7s 18ms/step - accuracy: 0.9886 - loss: 0.0359 - val_accuracy: 0.9904 - val_loss: 0.0339
 Epoch 9/10
 375/375 ————— 7s 17ms/step - accuracy: 0.9895 - loss: 0.0328 - val_accuracy: 0.9908 - val_loss: 0.0363
 Epoch 10/10
 375/375 ————— 7s 17ms/step - accuracy: 0.9916 - loss: 0.0261 - val_accuracy: 0.9908 - val_loss: 0.0343

Evaluate the Model

```
In [18]: # Step 5: Evaluate the Model
loss, accuracy = model.evaluate(x_test, y_test, verbose=0)
print(f"Test Accuracy: {accuracy * 100:.2f}%")
```


Test Accuracy: 99.19%

Confusion Matrix

```
In [ ]: # Confusion Matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)
conf_matrix = confusion_matrix(y_true, y_pred_classes)
```

Plot Confusion Matrix

```
In [23]: # Plot Confusion Matrix
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=range(10), yticklabels=range(10))
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



Classification Report

```
In [25]: # Classification Report
print(classification_report(y_true, y_pred_classes))
```

	precision	recall	f1-score	support
0	0.99	1.00	1.00	980
1	1.00	0.99	1.00	1135
2	1.00	0.99	0.99	1032
3	0.99	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.99	0.99	892
6	0.99	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

Plot Training History

```
In [27]: # Step 6: Plot Training History
plt.figure(figsize=(12, 4))
```

Out[27]: <Figure size 1200x400 with 0 Axes>


<Figure size 1200x400 with 0 Axes>

Accuracy Plot

```
In [29]: # Accuracy Plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
```

```
plt.legend()  
plt.title('Model Accuracy')
```


Out[29]: Text(0.5, 1.0, 'Model Accuracy')



Loss Plot

```
In [31]: # Loss Plot  
plt.subplot(1, 2, 2)  
plt.plot(history.history['loss'], label='Training Loss')  
plt.plot(history.history['val_loss'], label='Validation Loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()
```

```
plt.title('Model Loss')  
plt.show()
```



Save the Model

```
In [51]: # Save the Model  
model.save('mnist_cnn_model.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
In [ ]:
```